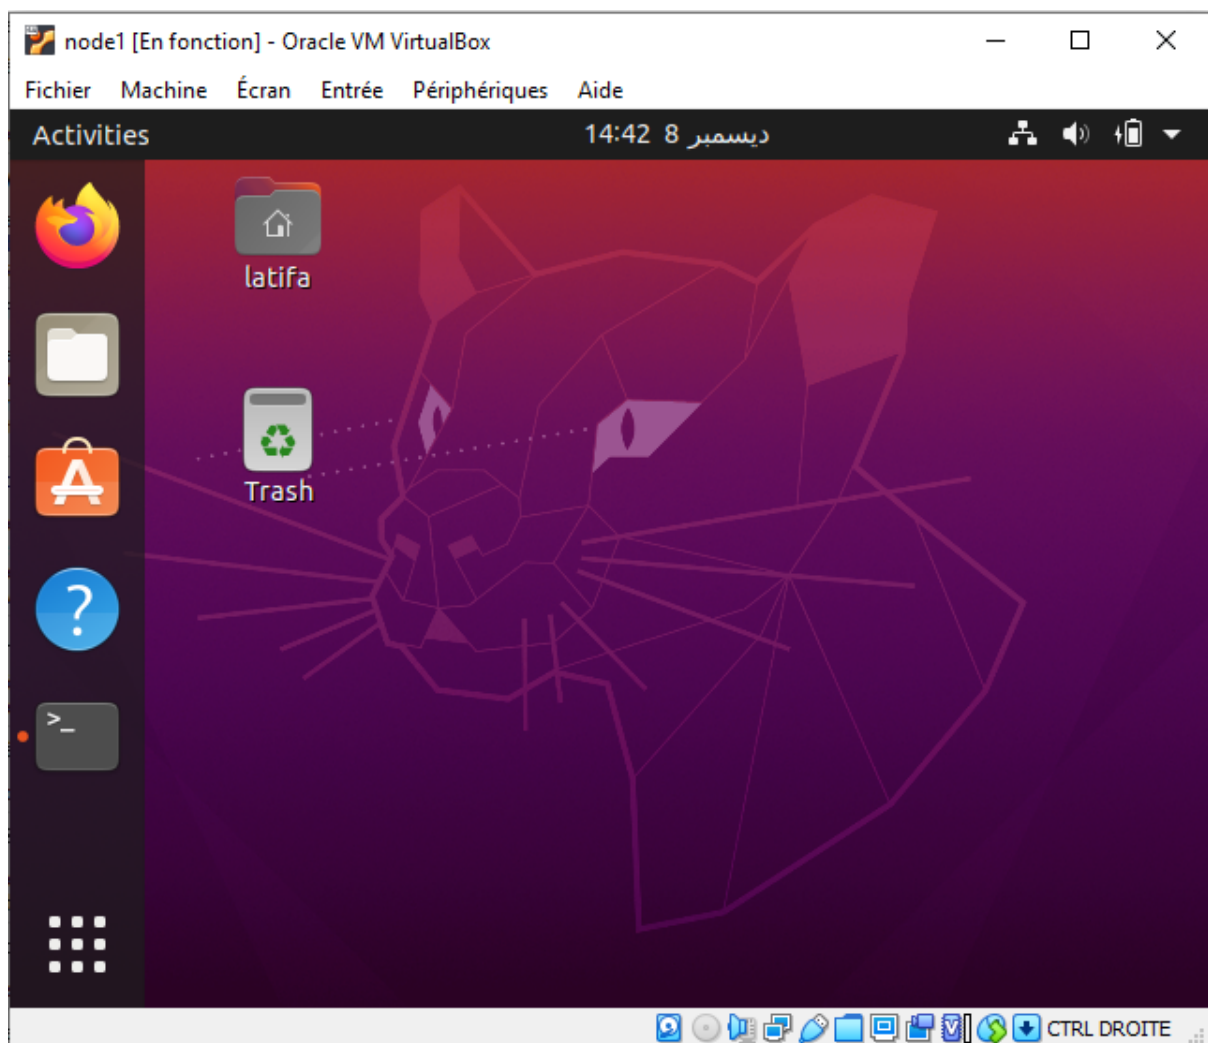


Projet Big Data

1. Création d'un cluster single node

a. Installation hadoop

La première étape pour installer hadoop est l'installation du JDK sur le système d'exploitation(dans notre cas ubuntu).



Tout d'abord il faut mettre à jour les fichiers du dépôt dans le système en tapant la commande:

```
latifa@latifa-VirtualBox:~$ sudo apt-get update
```

```
latifa@latifa-VirtualBox:~$ sudo apt-get update
[sudo] password for latifa:
Hit:1 http://tn.archive.ubuntu.com/ubuntu focal InRelease
Get:2 http://tn.archive.ubuntu.com/ubuntu focal-updates InRelease [114 kB]
Get:3 http://security.ubuntu.com/ubuntu focal-security InRelease [114 kB]
Get:4 http://tn.archive.ubuntu.com/ubuntu focal-backports InRelease [108 kB]
Get:5 http://tn.archive.ubuntu.com/ubuntu focal-updates/main amd64 Packages [1,391 kB]
Get:6 http://security.ubuntu.com/ubuntu focal-security/main i386 Packages [344 kB]
Get:7 http://tn.archive.ubuntu.com/ubuntu focal-updates/main i386 Packages [572
```

apt-get update : commande qui met à jour la liste des fichiers disponibles dans les dépôts APT présents dans le fichier de configuration /etc/apt/sources.list.

Ensuite, on installe le java JDK, dans notre projet, nous avons installé la version par défaut avec la commande:

```
latifa@latifa-VirtualBox:~$ sudo apt-get install default-jdk
Reading package lists... Done
Building dependency tree
Reading state information... Done
default-jdk is already the newest version (2:1.11-72).
0 upgraded, 0 newly installed, 0 to remove and 11 not upgraded.
```

Pour vérifier la version installée de java, on tape la commande:

```
latifa@latifa-VirtualBox:~$ java --version
openjdk 11.0.11 2021-04-20
OpenJDK Runtime Environment (build 11.0.11+9-Ubuntu-0ubuntu2.20.04)
OpenJDK 64-Bit Server VM (build 11.0.11+9-Ubuntu-0ubuntu2.20.04, mixed mode, sharing)
latifa@latifa-VirtualBox:~$
```

Il faut maintenant configurer l'environnement java avec l'installation du ssh:

```
latifa@latifa-VirtualBox:~$ sudo apt-get install ssh
Reading package lists... Done
Building dependency tree
Reading state information... Done
ssh is already the newest version (1:8.2p1-4ubuntu0.3).
```

```
latifa@latifa-VirtualBox:~$ sudo apt-get install rsync
```

rsync (pour la synchronisation à distance), est un logiciel de synchronisation de fichiers. Il est fréquemment utilisé pour mettre en place des systèmes de sauvegarde distante. On génère par la suite la clé du cryptage avec la commande:

```
latifa@latifa-VirtualBox:~$ ssh-keygen -t rsa -P ""
Generating public/private rsa key pair.
Enter file in which to save the key (/home/latifa/.ssh/id_rsa):
/home/latifa/.ssh/id_rsa already exists.
Overwrite (y/n)? y
Your identification has been saved in /home/latifa/.ssh/id_rsa
Your public key has been saved in /home/latifa/.ssh/id_rsa.pub
The key fingerprint is:
SHA256:QZCgBC5yEijMPA2Y2AS/KhRjUHNfknkFNkmRn3yUfpQ latifa@latifa-VirtualBox
The key's randomart image is:
+---[RSA 3072]-----+
|+*=.oB00o  . .|
|O+++oo=o  o E|
|=0=..  o.+ .|
|+++..  +.o .|
|+.+    S. .|
|o=|
|o|
|.|
|_|
+----[SHA256]-----+
latifa@latifa-VirtualBox:~$
```

Il faut copier la clé par la suite dans le répertoire Authorised keys avec la commande

```
|o|
|.|
|_|
+----[SHA256]-----+
latifa@latifa-VirtualBox:~$ cat $HOME/.ssh/id_rsa.pub >> $HOME/.ssh/authorized_k
keys
latifa@latifa-VirtualBox:~$
```

Pour vérifier l'installation, on tape la commande: ssh localhost

```
latifa@latifa-VirtualBox:~$ ssh localhost
Welcome to Ubuntu 20.04.3 LTS (GNU/Linux 5.11.0-41-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

1 update can be applied immediately.
To see these additional updates run: apt list --upgradable

Your Hardware Enablement Stack (HWE) is supported until April 2025.
Last login: Fri Dec  3 22:16:19 2021 from 127.0.0.1
latifa@latifa-VirtualBox:~$
```

Une fois qu'on a terminé la configuration du jdk, on passe maintenant à la désactivation du protocole IPV6, il faut donc taper la commande et ajouter les lignes ci-dessous dans le fichier sysctl.conf.

```
latifa@latifa-VirtualBox:~$ sudo gedit /etc/sysctl.conf
[sudo] password for latifa:
```

```
69 # disable ipv6
70 net.ipv6.conf.all.disable_ipv6 = 1
71 net.ipv6.conf.default.disable_ipv6 = 1
72 net.ipv6.conf.lo.disable_ipv6 = 1
73
```

Par la suite, on fait le téléchargement de la version d'Hadoop à partir du site officiel d'Apache ou simplement à partir du terminal et extraire le dossier zippé par la suite avec les commandes:

```
latifa@latifa-VirtualBox:~$ cd Downloads
latifa@latifa-VirtualBox:~/Downloads$ wget https://archive.apache.org/dist/hadoop/core/hadoop-2.6.0/hadoop-2.6.0.tar.gz
```

```
latifa@latifa-VirtualBox:~/Downloads$ sudo tar -xzf hadoop-2.6.0.tar.gz
```

On déplace le dossier dans le `usr/local/hadoop` et on attribue les droits d'accès au `user1` déjà créé avec la commande:

```
latifa@latifa-VirtualBox:~/Downloads$ sudo mv hadoop-2.6.0 /usr/local/hadoop
```

```
latifa@latifa-VirtualBox:~/Downloads$ ls /usr/local/hadoop
bin  include  libexec  logs      README.txt  share
etc  lib      LICENSE.txt  NOTICE.txt  sbin
```

```
latifa@latifa-VirtualBox:~/Downloads$ sudo chown user1:hadoop -R /usr/local/hadoop
latifa@latifa-VirtualBox:~/Downloads$
```

On crée par la suite les répertoires pour les datanode et le namenode et on applique les droits accès à ces deux répertoires aux `user1`.

```
user1@latifa-VirtualBox:~$ sudo mkdir -p /usr/local/hadoop_tmp/hdfs/namenode
```

```
user1@latifa-VirtualBox:~$ sudo mkdir -p /usr/local/hadoop_tmp/hdfs/datanode
```

```
user1@latifa-VirtualBox:~$ sudo chown user1:hadoop -R /usr/local/hadoop_tmp/
user1@latifa-VirtualBox:~$
```

b. Configuration Hadoop single node

Tout d'abord, il faut mettre à jour le fichier **bashrc** avec la commande:

```
user1@latifa-VirtualBox:~$ sudo gedit .bashrc
```

Et on ajoute les lignes au bas du fichier tout en spécifiant l'emplacement du JDK dans l'alias **JAVA_HOME** et le chemin d'hadoop dans l'alias **HADOOP_HOME**

```
fi
# -- HADOOP ENVIRONMENT VARIABLES START -- #
export JAVA_HOME=/usr/lib/jvm/java-11-openjdk-amd64
export HADOOP_HOME=/usr/local/hadoop
export PATH=$PATH:$HADOOP_HOME/bin
export PATH=$PATH:$HADOOP_HOME/sbin
export HADOOP_MAPRED_HOME=$HADOOP_HOME
export HADOOP_COMMON_HOME=$HADOOP_HOME
export HADOOP_HDFS_HOME=$HADOOP_HOME
export YARN_HOME=$HADOOP_HOME
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_HOME/lib/native
export HADOOP_OPTS="-Djava.library.path=$HADOOP_HOME/lib"
# -- HADOOP ENVIRONMENT VARIABLES END -- #
```

Pour que les alias puissent entrer en vigueur, **il faut redémarrer le terminal** ou taper la commande **source .bashrc**.

On peut maintenant faire la configuration de Hadoop en mode pseudo-distribué, il faut modifier 5 fichiers qui sont dans **/usr/local/hadoop/etc/hadoop/**

- **Mettre à jour le fichier env.sh**

Comme Hadoop est développé en Java, on doit lui spécifier le chemin de Jdk pour qu'il puisse activer ses démons.

```
user1@latifa-VirtualBox:~$ cd /usr/local/hadoop/etc/hadoop
user1@latifa-VirtualBox:/usr/local/hadoop/etc/hadoop$ sudo gedit hadoop-env.sh
[sudo] password for user1:
```

Il faut modifier la variable **JAVA_HOME** comme suit

```
100 |
101 export JAVA_HOME=/usr/lib/jvm/java-11-openjdk-amd64
```

- **Mettre à jour le fichier core-site.xml**

Le fichier core-site.xml informe les démons de hadoop qu'un namenode s'exécute sur le cluster en mentionnant son adresse.

Puisqu'on a une seule machine dans le cluster, le namenode va être sur le localhost (127.0.0.1). Le port 9000 est associé au système de fichier HDFS.

```
user1@latifa-VirtualBox:/usr/local/hadoop/etc/hadoop$ sudo gedit core-site.xml
```

Et on ajoute dans la balise configuration:

```
16
17 <!-- Put site-specific property overrides in this file. -->
18
19 <configuration>
20 <property>
21 <name>fs.default.name</name>
22 <value>hdfs://localhost:9000</value>
23 </property>
24 </configuration>
```

- **Mettre à jour le fichier hdfs-site.xml**

le fichier hdfs-site.xml informe hadoop et son système hdfs du nombre de réplifications (property 1), l'adresse de l'historique des transactions du NameNode (property 2) et l'adresse du stockage des blocs par les DataNode (property 3).

La réplication doit être 1 dans notre cas car on travaille sur un cluster à un seul noeud

On tape alors la commande et on ajoute les 3 propriétés dans la balise de configuration.

```
user1@latifa-VirtualBox:/usr/local/hadoop/etc/hadoop$ sudo gedit hdfs-site.xml
```

```
18
19 <configuration>
20 <property>
21 <name>dfs.replication</name>
22 <value>1</value>
23 </property>
24
25 <property>
26 <name>dfs.namenode.name.dir</name>
27 <value>file:/usr/local/hadoop_tmp/hdfs/namenode</value>
28 </property>
29
30 <property>
31 <name>dfs.datanode.data.dir</name>
32 <value>file:/usr/local/hadoop_tmp/hdfs/datanode</value>
33 </property>
34 </configuration>
```

XML ▼ Ta

- **Mettre à jour le fichier mapred-site.xml**

Ce fichier mapred-site.xml informe le package MapReduce qu'il va s'exécuter en tant qu'une application yarn (séparation entre la gestion des ressources et la gestion des

traitements). Avant de le modifier, il faut d'abord faire une copie du fichier `mapred-site.xml.template` sous le nom `mapred-site.xml` avec la commande :

```
user1@latifa-VirtualBox:/usr/local/hadoop/etc/hadoop$ sudo cp mapred-site.xml.template mapred-site.xml
```

```
user1@latifa-VirtualBox:/usr/local/hadoop/etc/hadoop$ sudo gedit mapred-site.xml
```

Dans la balise `configuration`, on ajoute les lignes suivantes :

```
18
19 <configuration>
20 <property>
21 <name>mapreduce.framework.name</name>
22 <value>yarn</value>
23 </property>
24
25 </configuration>
```

- **Mettre à jour le fichier `yarn-site.xml`**

Le fichier `yarn-site.xml` indique au node manager qu'il aura un service auxiliaire indiquant au MapReduce comment faire son shuffling.

```
user1@latifa-VirtualBox:/usr/local/hadoop/etc/hadoop$ sudo gedit yarn-site.xml
```

```
15 <configuration>
16 <property>
17 <name>yarn.nodemanager.aux-services</name>
18 <value>mapreduce_shuffle</value>
19 </property>
20
21 <property>
22 <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
23 <value>org.apache.hadoop.mapred.ShuffleHandler</value>
24 </property>
25 <!-- Site specific YARN configuration properties -->
26
27 </configuration>
```

c. Vérifier l'installation

Une fois qu'on a terminé la configuration de hadoop, on va vérifier son installation. Pour ce faire, il faut tout d'abord formater le NameNode avant de démarrer Hadoop et ça se fait à chaque fois qu'on va démarrer les services de Hadoop.

```
user1@latifa-VirtualBox:/usr/local/hadoop/etc/hadoop$ hdfs namenode -format
21/12/08 17:22:32 INFO namenode.NameNode: STARTUP_MSG:
/*****
STARTUP_MSG: Starting NameNode
STARTUP_MSG:   host = latifa-VirtualBox/127.0.1.1
STARTUP_MSG:   args = [-format]
STARTUP_MSG:   version = 2.6.0
STARTUP_MSG:   classpath = /usr/local/hadoop/etc/hadoop:/usr/local/hadoop/share/
hadoop/common/lib/curator-client-2.6.0.jar:/usr/local/hadoop/share/hadoop/common
/lib/jackson-jaxrs-1.9.13.jar:/usr/local/hadoop/share/hadoop/common/lib/log4j-1.
2.17.jar:/usr/local/hadoop/share/hadoop/common/lib/jsch-0.1.42.jar:/usr/local/ha
doop/share/hadoop/common/lib/jasper-runtime-5.5.23.jar:/usr/local/hadoop/share/h
adoop/common/lib/jersey-core-1.9.jar:/usr/local/hadoop/share/hadoop/common/lib/j
sp-api-2.1.jar:/usr/local/hadoop/share/hadoop/common/lib/slf4j-api-1.7.5.jar:/us
r/local/hadoop/share/hadoop/common/lib/commons-cli-1.2.jar:/usr/local/hadoop/sha
re/hadoop/common/lib/gson-2.2.4.jar:/usr/local/hadoop/share/hadoop/common/lib/ht
tpcore-4.2.5.jar:/usr/local/hadoop/share/hadoop/common/lib/apacheds-kerberos-cod
ec-2.0.0-M15.jar:/usr/local/hadoop/share/hadoop/common/lib/avro-1.7.4.jar:/usr/l
ocal/hadoop/share/hadoop/common/lib/protobuf-java-2.5.0.jar:/usr/local/hadoop/sh
```

d. Démarrer le processus Hadoop

```
user1@latifa-VirtualBox:/usr/local/hadoop/etc/hadoop$ start-dfs.sh
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.
util.KerberosUtil (file:/usr/local/hadoop/share/hadoop/common/lib/hadoop-auth-2.
6.0.jar) to method sun.security.krb5.Config.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.
security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflect
```

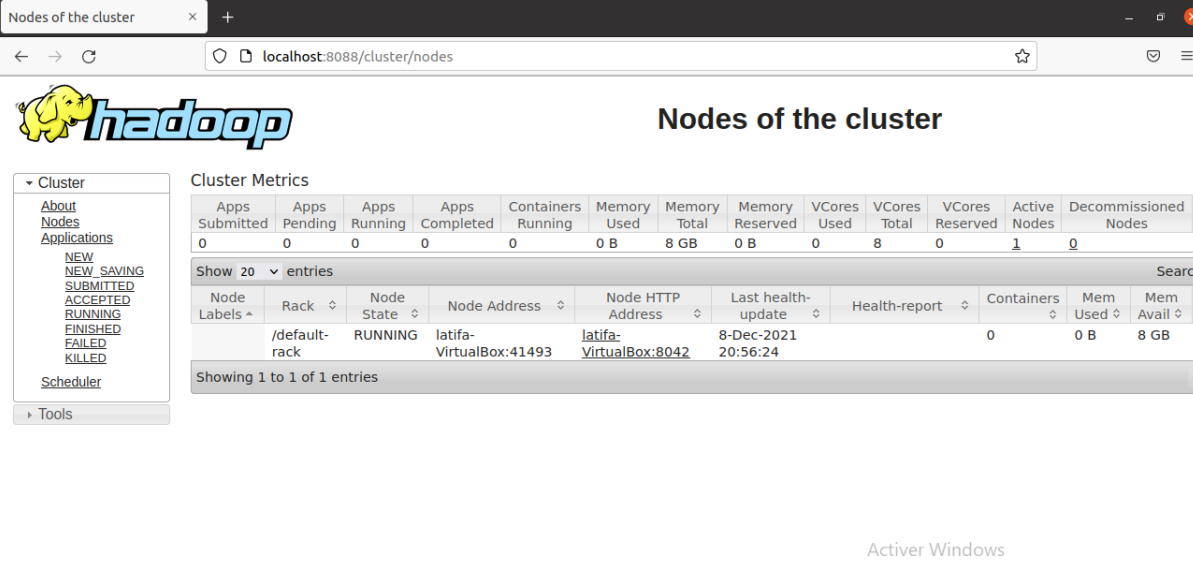
```
user1@latifa-VirtualBox:/usr/local/hadoop/etc/hadoop$ start-yarn.sh
starting yarn daemons
starting resourcemanager, logging to /usr/local/hadoop/logs/yarn-user1-resourcem
anager-latifa-VirtualBox.out
localhost: starting nodemanager, logging to /usr/local/hadoop/logs/yarn-user1-no
demanager-latifa-VirtualBox.out
user1@latifa-VirtualBox:/usr/local/hadoop/etc/hadoop$
```

e. Vérifier le service après le démarrage

```
user1@latifa-VirtualBox:/usr/local/hadoop/etc/hadoop$ jps
11649 ResourceManager
11080 NameNode
12073 Jps
11454 SecondaryNameNode
11775 NodeManager
user1@latifa-VirtualBox:/usr/local/hadoop/etc/hadoop$
```


f. Visualisation de l'interface graphique du NameNode

Quand on se connecte sur <http://localhost:8080> une interface se présente indiquant un seul noeuds actif.



The screenshot shows the Hadoop web interface for the NameNode. The browser address bar displays `localhost:8088/cluster/nodes`. The page title is "Nodes of the cluster". On the left, there is a sidebar with a "Cluster" menu containing links for "About", "Nodes", "Applications", "NEW", "NEW SAVING", "SUBMITTED", "ACCEPTED", "RUNNING", "FINISHED", "FAILED", "KILLED", "Scheduler", and "Tools". The main content area displays "Cluster Metrics" with a table showing various metrics. Below this, there is a table of nodes with columns for Node Labels, Rack, Node State, Node Address, Node HTTP Address, Last health-update, Health-report, Containers, Mem Used, and Mem Avail. The table shows one node in the "RUNNING" state with the address "latifa-VirtualBox:41493" and "latifa-VirtualBox:8042".

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved	Active Nodes	Decommissioned Nodes
0	0	0	0	0	0 B	8 GB	0 B	0	8	0	1	0

Node Labels	Rack	Node State	Node Address	Node HTTP Address	Last health-update	Health-report	Containers	Mem Used	Mem Avail
/default-rack		RUNNING	latifa-VirtualBox:41493	latifa-VirtualBox:8042	8-Dec-2021 20:56:24		0	0 B	8 GB



Overview 'localhost:9000' (active)

Started:	Wed Dec 08 17:28:13 CET 2021
Version:	2.6.0, re3496499ecb8d220fba99dc5ed4c99c8f9e33bb1
Compiled:	2014-11-13T21:10Z by jenkins from (detached from e349649)
Cluster ID:	CID-dc97fd56-ee18-4582-b393-a252d6b33173
Block Pool ID:	BP-1954933342-127.0.1.1-1638980558142

Summary

Security is off.

Activier Windows

2. Création d'un cluster multi-nodes

On a choisi de travailler avec un cluster à 3 nœuds: Un NameNode et deux DataNodes. Par conséquent, on doit modifier la configuration précédente pour permettre d'exécuter les *jobs* de façon distribuée.

La première chose à faire est de dupliquer la machine qu'on a déjà créé en mode pseudo-distribué et mettre la configuration réseau en mode accès par pont afin d'avoir des adresses IP différentes dans les trois machines.

Après la duplication, il faut obtenir les adresses IP des trois machines virtuelles.

```
latifa@NameNode:~$ ip addr show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:1d:49:1c brd ff:ff:ff:ff:ff:ff
    inet 192.168.100.131/24 brd 192.168.100.255 scope global dynamic noprefixroute enp0s3
        valid_lft 85758sec preferred_lft 85758sec
    inet6 fe80::32c9:c6c8:4478:f686/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
```

```
latifa@DataNode1:~$ ip addr show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:12:1e:51 brd ff:ff:ff:ff:ff:ff
    inet 192.168.100.168/24 brd 192.168.100.255 scope global dynamic noprefixroute enp0s3
        valid_lft 85659sec preferred_lft 85659sec
    inet6 fe80::f0b6:19f3:dd7b:3d6f/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
```

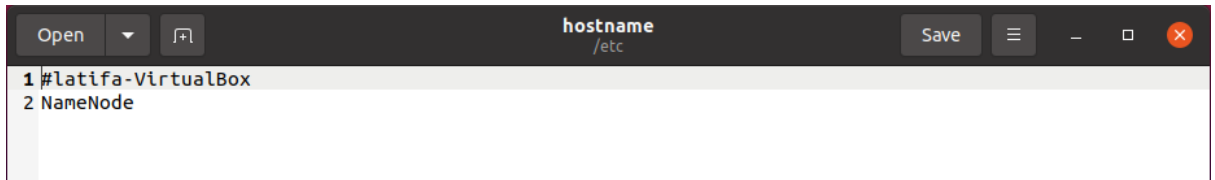
```
latifa@DataNode2:~$ ip addr show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:41:86:4b brd ff:ff:ff:ff:ff:ff
    inet 192.168.100.183/24 brd 192.168.100.255 scope global dynamic noprefixroute enp0s3
        valid_lft 85631sec preferred_lft 85631sec
    inet6 fe80::bfe4:62c0:61a:73e8/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
```

a. Configuration des alias réseaux:

On édite le fichier /etc/hostname en tant qu'utilisateur root dans les trois Vms.

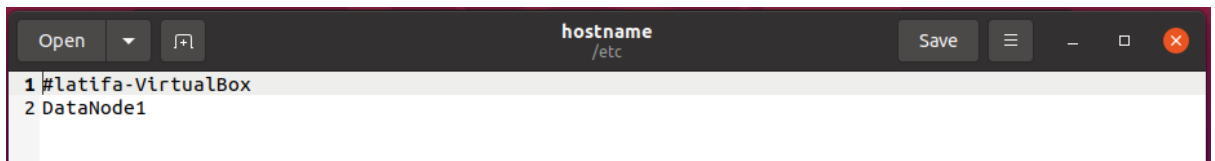
```
latifa@DataNode1:~$ su
Password:
root@DataNode1:/home/latifa# gedit /etc/hostname
```

- Dans le nameNode:



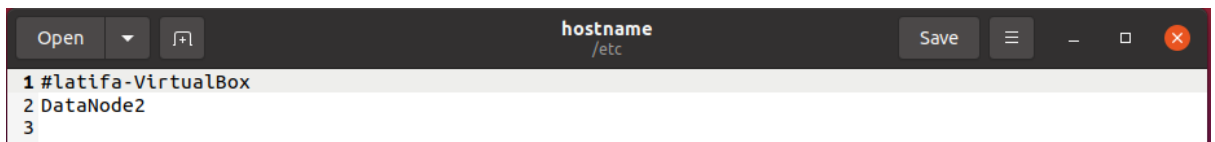
A screenshot of a text editor window titled 'hostname /etc'. The window contains two lines of text: '1 #latifa-VirtualBox' and '2 NameNode'.

- Dans le DataNode 1



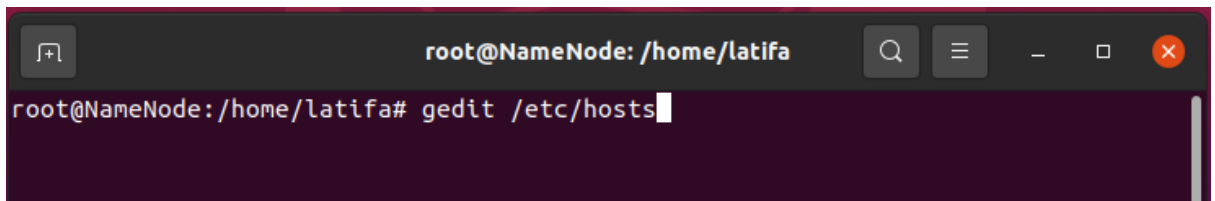
A screenshot of a text editor window titled 'hostname /etc'. The window contains two lines of text: '1 #latifa-VirtualBox' and '2 DataNode1'.

- Dans le DataNode 2

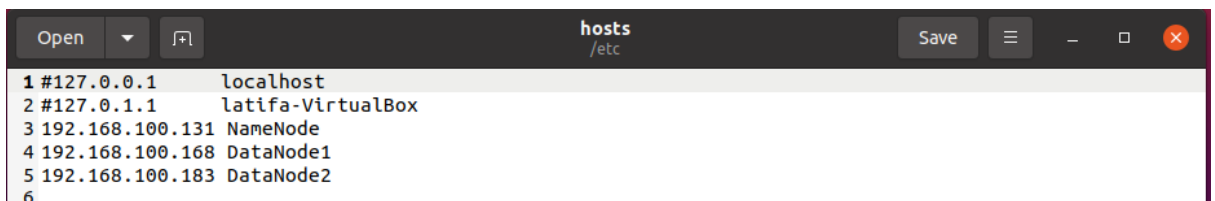


A screenshot of a text editor window titled 'hostname /etc'. The window contains three lines of text: '1 #latifa-VirtualBox', '2 DataNode2', and '3'.

On édite aussi le fichier /etc/hosts dans les 3 machines virtuelles comme suit:



A screenshot of a terminal window showing the command 'root@NameNode:/home/latifa# gedit /etc/hosts' being entered.



A screenshot of a text editor window titled 'hosts /etc'. The window contains six lines of text: '1 #127.0.0.1 localhost', '2 #127.0.0.1 latifa-VirtualBox', '3 192.168.100.131 NameNode', '4 192.168.100.168 DataNode1', '5 192.168.100.183 DataNode2', and '6'.

Après cette configuration, il faut redémarrer les VMs avec la commande reboot.

b. Configuration des fichiers Xml:

Ces configurations sont communes aux trois nœuds, on doit donc se connecter en tant qu'utilisateur user1 déjà créé.

- Configuration du fichier core-site.xml:

```
user1@NameNode:~$ su user1
Password:
user1@NameNode:~$ cd /usr/local/hadoop/etc/hadoop
user1@NameNode:/usr/local/hadoop/etc/hadoop$ sudo gedit core-site.xml
```

On remplace la valeur de la propriété `fs.default.name` `hdfs://localhost:9000` par `hdfs://NameNode:9000` puisque l'alias réseau a été modifié dans le fichier `/etc/hosts`.

```
16
17 <!-- Put site-specific property overrides in this file. -->
18
19 <configuration>
20 <property>
21 <name>fs.default.name</name>
22 <value>hdfs://NameNode:9000</value>
23 </property>
24 </configuration>
```

XML ▾ Tab Width: 8 ▾ Ln 21, Col 29 ▾ INS

- **Configuration du fichier `hdfs-site.xml`:**

```
(gedit:2270): Tepl-WARNING **: 00:18:11.342: GVfs metadata is not supported. Fal
lback to TeplMetadataManager. Either GVfs is not correctly installed or GVfs met
adata are not supported on this platform. In the latter case, you should configu
re Tepl with --disable-gvfs-metadata.
user1@NameNode:/usr/local/hadoop/etc/hadoop$ sudo gedit hdfs-site.xml
```

On change le facteur de réplication à 2 dans la propriété `dfs.replication`.

```
19 <configuration>
20 <property>
21 <name>dfs.replication</name>
22 <value>2</value>
23 </property>
24
25 <property>
26 <name>dfs.namenode.name.dir</name>
27 <value>file:/usr/local/hadoop_tmp/hdfs/namenode</value>
28 </property>
29
30 <property>
31 <name>dfs.datanode.data.dir</name>
32 <value>file:/usr/local/hadoop_tmp/hdfs/datanode</value>
33 </property>
34 </configuration>
```

- **Configuration du fichier `yarn-site.xml`:**

```
user1@NameNode: /usr/local/hadoop/etc/hadoop
user1@NameNode: /usr/local/hadoop/etc/hadoop$ sudo gedit yarn-site.xml
```

```
14 -->
15 <configuration>
16 <property>
17 <name>yarn.nodemanager.aux-services</name>
18 <value>mapreduce_shuffle</value>
19 </property>
20 <property>
21 <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
22 <value>org.apache.hadoop.mapred.ShuffleHandler</value>
23 </property>
24 <property>
25 <name>yarn.resourcemanager.resource-tracker.address</name>
26 <value>NameNode:8025</value>
27 </property>
28
29 <property>
30 <name>yarn.resourcemanager.scheduler.address</name>
31 <value>NameNode:8035</value>
32 </property>
33
34 <property>
35 <name>yarn.resourcemanager.address</name>
36 <value>NameNode:8050</value>
37 </property>
38
39 </configuration>
```

- **Configuration du fichier mapred-site.xml:**

```
(gedit:2302): Tepl-WARNING **: 00:26:37.584: GVfs metadata is not supported. Fall
back to TeplMetadataManager. Either GVfs is not correctly installed or GVfs met
adata are not supported on this platform. In the latter case, you should configu
re Tepl with --disable-gvfs-metadata.
user1@NameNode: /usr/local/hadoop/etc/hadoop$ sudo gedit mapred-site.xml
```

On modifie la valeur du paramètre `mapreduce.framework.name` en la passant à `yarn`.

```
18
19 <configuration>
20 <property>
21 <name>mapreduce.job.tracker</name>
22 <value>NameNode:5431</value>
23 </property>
24 <property>
25 <name>mapreduce.framework.name</name>
26 <value>yarn</value>
27 </property>
28
29 </configuration>
```

c. Configuration des fichiers masters et slaves:

```
user1@NameNode: /usr/local/hadoop/etc/hadoop
user1@NameNode:/usr/local/hadoop/etc/hadoop$ sudo gedit masters
[sudo] password for user1:
```

On spécifie dans le fichier le nom du master.

```
Open  ▾  +  masters /usr/local/hadoop/etc/hadoop Save  ≡  -  □  ×
1 NameNode
```

```
user1@NameNode: /usr/local/hadoop/etc/hadoop
user1@NameNode:/usr/local/hadoop/etc/hadoop$ sudo gedit slaves
```

On spécifie dans le fichier le nom des workers.

```
Open  ▾  +  slaves /usr/local/hadoop/etc/hadoop Save  ≡  -  □  ×
1 #localhost
2 DataNode1
3 DataNode2
4
```

d. Formatage du système HDFS dans le NameNode:

On supprime le dossier **hadoop_tmp** existant et on crée lors de la configuration pseudo-distribué.


```
user1@NameNode: ~  
user1@NameNode:~$ sudo rm -rf /usr/local/hadoop_tmp/
```

Et on recrée par la suite le répertoire NameNode et on donne les droits d'accès de ce répertoire au user1.

```
user1@NameNode:~$ sudo mkdir -p /usr/local/hadoop_tmp/
```

```
user1@NameNode:~$ sudo mkdir -p /usr/local/hadoop_tmp/hdfs/namenode
```

```
user1@NameNode:~$ sudo chown user1:hadoop -R /usr/local/hadoop_tmp/
```

e. Démarrage des services:

```
user1@NameNode:~$ hdfs namenode -format  
21/12/11 17:41:11 INFO namenode.NameNode: STARTUP_MSG:  
/*****  
STARTUP_MSG: Starting NameNode  
STARTUP_MSG:   host = NameNode/192.168.100.120  
STARTUP_MSG:   args = [-format]  
STARTUP_MSG:   version = 2.6.0  
STARTUP_MSG:   classpath = /usr/local/hadoop/etc/hadoop:/usr/local/hadoop/share/  
hadoop/common/lib/curator-client-2.6.0.jar:/usr/local/hadoop/share/hadoop/common  
/lib/jackson-jaxrs-1.9.13.jar:/usr/local/hadoop/share/hadoop/common/lib/log4j-1.  
2.17.jar:/usr/local/hadoop/share/hadoop/common/lib/jsch-0.1.42.jar:/usr/local/ha
```

```
user1@NameNode:~$ start-dfs.sh  
WARNING: An illegal reflective access operation has occurred  
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.  
util.KerberosUtil (file:/usr/local/hadoop/share/hadoop/common/lib/hadoop-auth-2.  
6.0.jar) to method sun.security.krb5.Config.getInstance()  
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.  
security.authentication.util.KerberosUtil  
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflect  
ive access operations  
WARNING: All illegal access operations will be denied in a future release
```

```
user1@NameNode:~$ start-yarn.sh
starting yarn daemons
starting resourcemanager, logging to /usr/local/hadoop/logs/yarn-user1-resourcemanager-NameNode.out
DataNode2: starting nodemanager, logging to /usr/local/hadoop/logs/yarn-user1-nodemanager-DataNode2.out
DataNode1: starting nodemanager, logging to /usr/local/hadoop/logs/yarn-user1-nodemanager-DataNode1.out
```

Résultat dans le Master

```
user1@NameNode:/usr/local/hadoop/etc/hadoop$ jps
3712 DataNode
4065 NodeManager
2068 SecondaryNameNode
1831 NameNode
4093 Jps
2207 ResourceManager
user1@NameNode:/usr/local/hadoop/etc/hadoop$
```

Résultat dans les workers

```
user1@DataNode2:~$ jps
2625 NodeManager
2417 DataNode
2852 Jps
user1@DataNode2:~$
```

```
user1@DataNode1:~$ jps
2452 DataNode
2888 Jps
2664 NodeManager
user1@DataNode1:~$
```

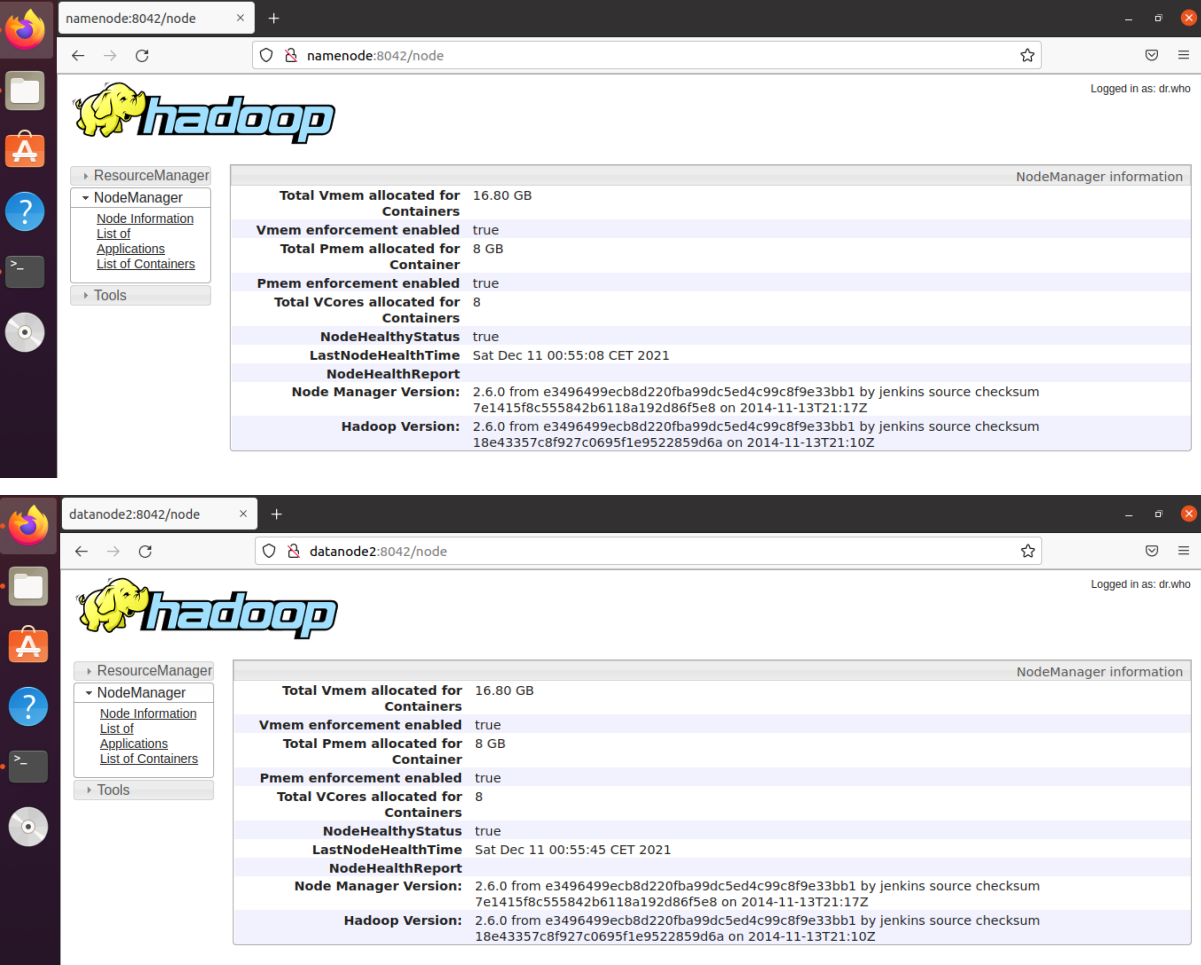
f. Visualisation des résultats:

The screenshot shows the Hadoop web interface in a Firefox browser window. The page title is 'Nodes of the cluster'. On the left, there is a sidebar with a 'Cluster' section containing links for 'About', 'Nodes', and 'Applications'. The 'Nodes' link is selected. The main content area shows 'Cluster Metrics' and a table of nodes.

Cluster Metrics											
Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved	Active Nodes
0	0	0	0	0	0 B	24 GB	0 B	0	24	0	3

Node Labels	Rack	Node State	Node Address	Node HTTP Address	Last health-update	Health-report	Containers	Mem Used	Mem Avail
/default-rack		RUNNING	NameNode:44013	NameNode:8042	11-Dec-2021 00:49:07		0	0 B	8 GB
/default-rack		RUNNING	DataNode2:38899	DataNode2:8042	11-Dec-2021 00:47:45		0	0 B	8 GB
/default-rack		RUNNING	DataNode1:40217	DataNode1:8042	11-Dec-2021 00:49:03		0	0 B	8 GB

Showing 1 to 3 of 3 entries



The image displays two screenshots of the Hadoop web interface, specifically the NodeManager information page. Both screenshots show the same information, indicating a consistent configuration across the nodes.

NodeManager information

Total Vmem allocated for Containers	16.80 GB
Vmem enforcement enabled	true
Total Pmem allocated for Container	8 GB
Pmem enforcement enabled	true
Total VCoers allocated for Containers	8
NodeHealthyStatus	true
LastNodeHealthTime	Sat Dec 11 00:55:08 CET 2021
NodeHealthReport	
Node Manager Version:	2.6.0 from e3496499ecb8d220fba99dc5ed4c99c8f9e33bb1 by jenkins source checksum 7e1415f8c555842b6118a192d86f5e8 on 2014-11-13T21:17Z
Hadoop Version:	2.6.0 from e3496499ecb8d220fba99dc5ed4c99c8f9e33bb1 by jenkins source checksum 18e43357c8f927c0695f1e9522859d6a on 2014-11-13T21:10Z

g. Exécution d'un Jar:

```
user1@NameNode:~$ hdfs dfs -mkdir /test/
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.
util.KerberosUtil (file:/usr/local/hadoop/share/hadoop/common/lib/hadoop-auth-2.
6.0.jar) to method sun.security.krb5.Config.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.
security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflect
ive access operations
WARNING: All illegal access operations will be denied in a future release
21/12/14 00:10:53 WARN util.NativeCodeLoader: Unable to load native-hadoop libra
ry for your platform... using builtin-java classes where applicable
mkdir: `/test': File exists
```

```
user1@NameNode:~$ hdfs dfs -put /home/latifa/wordCountinput.txt.txt /test/
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.
util.KerberosUtil (file:/usr/local/hadoop/share/hadoop/common/lib/hadoop-auth-2.
6.0.jar) to method sun.security.krb5.Config.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.
security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflect
ive access operations
WARNING: All illegal access operations will be denied in a future release
21/12/14 00:14:40 WARN util.NativeCodeLoader: Unable to load native-hadoop libra
ry for your platform... using builtin-java classes where applicable
```

```
user1@NameNode:~$ hadoop jar /usr/local/hadoop/share/hadoop/mapreduce/hadoop-map
reduce-examples-2.6.0.jar wordcount /test/wordCountinput.txt.txt /test/output/
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.
util.KerberosUtil (file:/usr/local/hadoop/share/hadoop/common/lib/hadoop-auth-2.
```

```
21/12/14 00:32:03 INFO mapreduce.Job: Job job_1639436260463_0002 completed succ
ssfully
```

```
21/12/14 00:32:16 INFO mapreduce.Job: Counters: 49
```

File System Counters

```
FILE: Number of bytes read=469
FILE: Number of bytes written=212627
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=392
HDFS: Number of bytes written=311
HDFS: Number of read operations=6
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
```

Job Counters

```
Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=32312
Total time spent by all reduces in occupied slots (ms)=264820
Total time spent by all map tasks (ms)=32312
Total time spent by all reduce tasks (ms)=264820
Total vcore-seconds taken by all map tasks=32312
Total vcore-seconds taken by all reduce tasks=264820
```